

# Project 2 Write-up

This is the write-up for project 2. We answer each question (to our best ability, shortcomings are mentioned) below.

## 1

### Comparison Table

Overall Report						
Metric/Code Version	Serial	Multicore	GPU1	GPU2	GPU3	GPU4
Time to solution (s)	924.647004	40.634892	4.324362	2.835888	2.281643	2.259270
Grind rate (s)	48.665632	2.138679	0.227598	0.149257	0.120086	0.118909
GFLOP/s	4.109	92.49505				
% time in force computation	99.9996184489881	99.9662457574638	99.8319983387145	99.7098616024328	99.7125317150843	99.6081477645434
Peak Theoretical (GLOP/s)	4.8	134.4				
% of peak	0.856041666666667	0.688207217261905				

Figure 1: Overall Report

Note: I had trouble getting metrics to work in order to calculating FLOP/s numerically on GPU. Hence the blank section.

We can clearly see that this code benefits greatly when run on a GPU (which makes sense given its SIMD nature)

## 2

### Serial Time Complexity

The complexity can be outlined by the following table.

The last column shows the relative ratio of time to nbody. Given the terms, 9, 18, 27, 37... we can roughly say that this routine has  $O(n^2)$  complexity.

This makes sense, given from analytic perspective, we are calculating forces between each pair of particles, which should naturally arrive at a  $O(n^2)$  runtime assuming a serial computer.

## 3

### Analytic Arithmetic Complexity

We can compute the FLOP/s for our serial implementation loop as such




Figure 2: Serial Scaling

Position Update: The position update is straightforward, it contains 3 adds and 3 multiplies on single-precision floats. This is done over N

BodyForce Calculation: - For each N, we do 3 subtractions, 9 mults, 6 adds, 1 div over N. Not sure how many floating point operations are in sqrtf() - At the end, we update our particles' velocities. via 3 adds and 3 multiplies. This is done over N.

So for each word, we have the following FLOP count. - 6 ops for pos update. -  $N \cdot 19$  ops for force calculation. - 6 ops for velocity update.

Yields  $12 + 19n$  FLOPs per word.

For  $n=100000$  -

$$100000 * (12 + 19 \cdot 100000) = 190001200000 \text{ FLOPs}$$

$$190001200000 / 1000000000 / 924.647 = 0.20548512026$$

Now clearly 0.20548512026 FLOP/s is nowhere close to our measured 4.109 FLOP/s. We ignored the sqrtf() function call, which I imagine contains more than 1 floating point operation. The discrepancy would imply that there are 20 such operations in sqrtf(), which isn't unimaginable. Kinda wish I had more insight into this than this (this is the last question I'm doing at the moment, I'm tired to be honest :) ).

Another thing to look at would be the re-run our above calculations with ffast-math turned on so that it yields a single FLOP for sqrt operations. We could then better directly compare the numerical results that way to verify our paper and pencil approach here.

## 4

### GPU v. CPU Speedup

Referencing our table in #1, we can see that the peak CPU scaling is at 40s and the peak GPU scaling is at 2.25s. This nets us a relative speedup of 17.98.

(It should be pretty clear that GPU will definitely have a higher performance/dollar – considering GPU cost scalability is definitely not north of 18)

## 5

### Amdahl Ceiling n=100000

Using a thread count of 32 and comparing against our serial run...

We have the following equation, where we can solve for the fraction of execution time affected by the speedup. We can use this to get our percentage of sequential execution (and optimal bound).

$$speedup = (924.647/40.63) = 22.757$$

$$22.757 = \frac{1}{(1 - p) + (p/32)}$$

Solving for p gives us 0.9868. Which tells us that our bound (our sequential) is at 1.3%.

## 6

### Amdahl Ceiling n=500000

Using a thread count of 32...following the same steps as above.

Sequential 500k = ? Parallel32 500k = 1016.639743

Note: Doing a sequential run for 500k will probably take more time than I can afford right now. Will come back around if I have time remaining.

## 7

### Compute or Memory Bound

My approach here given the relatively limited set of counters we have is to compare of ratio of the following two counters.

- PAPI\_STL\_CCY (Papi cycles that resulted in no instructions being completed)
- PAPI\_TOT\_CYC (Papi total cycles)

The thought here is that a higher ratio of PAPI\_STL\_CCY / PAPI\_TOT\_CYC would indicate a memory bound routine, a lesser ratio indicating a cpu bound routine.

When running for n=50000...

PAPI\_STL\_CCY = 273429250826

PAPI\_TOT\_CYC = 800989736575

Which gives us roughly 34% stalled cycles. This would indicate that our code is still compute-bound and could see further improvement if clock speed were increased (with a 0.66 improvement, essentially)

## 8

### CPU ICC

Using the timing from the next question (#9). We compare the following

- GCC: 924.647 (from #1)
- ICC: 265.059 (from #9, with vectorization)

Actually if you look at the no-vec from #9, we can see that the vast majority of this speedup is due to ICC vectorization in O2. I imagine in GCC, vectorization isn't done by default at O2.

## 9

### Vectorized Instructions

The following loops/instructions get vectorized in ICC (made sure to turn on AVX vectorization explicitly as well)

- Inner loop (only) gets vectorized in our force calculation function.
- The position update loop gets fully vectorized.

When comparing vectorization against no vec - W/vec: 265.058975 - Wo/vec: 923.806719

So clearly we see significant improvement when explicitly allowing vectorization to occur. This makes sense since the major computationally intensive loop in our

force computation subroutine is naturally a SIMD operation, which will benefit vastly from vectorization in ICC.

## Vectorization Report

Intel(R) Advisor can now assist with vectorization and show optimization report messages with your source code.

See "<https://software.intel.com/en-us/intel-advisor-xe>" for details.

Report from: Interprocedural optimizations [ipo]

INLINING OPTION VALUES:

-inline-factor: 100  
-inline-min-size: 30  
-inline-max-size: 230  
-inline-max-total-size: 2000  
-inline-max-per-routine: 10000  
-inline-max-per-compile: 500000

Begin optimization report for: main(const int, const char \*\*)

Report from: Interprocedural optimizations [ipo]

INLINE REPORT: (main(const int, const char \*\*)) [1] /home/jab29/scratch-midway2/project-2/src  
-> INLINE: (44,27) atoi(const char \*)  
{ { Inlining of routines from system headers is omitted. Use -qopt-report=3 to view full report }  
-> INLINE: (45,28) atoi(const char \*)  
{ { Inlining of routines from system headers is omitted. Use -qopt-report=3 to view full report }  
-> INLINE: (54,3) randomizeBodies(float \*, int)  
-> INLINE: (82,5) bodyForce(Body \*, float, int)

Report from: Loop nest, Vector & Auto-parallelization optimizations [loop, vec, par]

LOOP BEGIN at /home/jab29/scratch-midway2/project-2/src/nbody\_cpu\_serial.c(15,3) inlined int  
remark #15527: loop was not vectorized: function call to rand(void) cannot be vectorized  
LOOP END

LOOP BEGIN at /home/jab29/scratch-midway2/project-2/src/nbody\_cpu\_serial.c(21,3) inlined int  
remark #25236: Loop with pragma of trip count = 3000 ignored for large value  
remark #15542: loop was not vectorized: inner loop was already vectorized

LOOP BEGIN at /home/jab29/scratch-midway2/project-2/src/nbody\_cpu\_serial.c(24,5) inlined int

remark #15300: LOOP WAS VECTORIZED  
LOOP END

LOOP BEGIN at /home/jab29/scratch-midway2/project-2/src/nbody\_cpu\_serial.c(24,5) inlined int  
<Remainder loop for vectorization>  
LOOP END  
LOOP END

LOOP BEGIN at /home/jab29/scratch-midway2/project-2/src/nbody\_cpu\_serial.c(85,5)  
remark #15300: LOOP WAS VECTORIZED  
LOOP END

LOOP BEGIN at /home/jab29/scratch-midway2/project-2/src/nbody\_cpu\_serial.c(85,5)  
<Remainder loop for vectorization>  
remark #15301: REMAINDER LOOP WAS VECTORIZED  
LOOP END

LOOP BEGIN at /home/jab29/scratch-midway2/project-2/src/nbody\_cpu\_serial.c(85,5)  
<Remainder loop for vectorization>  
LOOP END

LOOP BEGIN at /home/jab29/scratch-midway2/project-2/src/nbody\_cpu\_serial.c(100,3)  
remark #15344: loop was not vectorized: vector dependence prevents vectorization. First dep  
LOOP END

Non-optimizable loops:

LOOP BEGIN at /home/jab29/scratch-midway2/project-2/src/nbody\_cpu\_serial.c(76,3)  
remark #15543: loop was not vectorized: loop with function call not considered an optimization  
LOOP END

Report from: Code generation optimizations [cg]

/home/jab29/scratch-midway2/project-2/src/nbody\_cpu\_serial.c(25,18):remark #34032: adjacent  
/home/jab29/scratch-midway2/project-2/src/nbody\_cpu\_serial.c(86,7):remark #34032: adjacent s  
/home/jab29/scratch-midway2/project-2/src/nbody\_cpu\_serial.c(86,7):remark #34032: adjacent s  
/home/jab29/scratch-midway2/project-2/src/nbody\_cpu\_serial.c(88,7):remark #34033: adjacent s  
=====

Begin optimization report for: randomizeBodies(float \*, int)

Report from: Interprocedural optimizations [ipo]

INLINE REPORT: (randomizeBodies(float \*, int)) [3] /home/jab29/scratch-midway2/project-2/src

Report from: Loop nest, Vector & Auto-parallelization optimizations [loop, vec, par]

```
LOOP BEGIN at /home/jab29/scratch-midway2/project-2/src/nbody_cpu_serial.c(15,3)
remark #15527: loop was not vectorized: function call to rand(void) cannot be vectorized
LOOP END
=====
```

Begin optimization report for: bodyForce(Body \*, float, int)

Report from: Interprocedural optimizations [ipo]

INLINE REPORT: (bodyForce(Body \*, float, int)) [4] /home/jab29/scratch-midway2/project-2/src

Report from: Loop nest, Vector & Auto-parallelization optimizations [loop, vec, par]

```
LOOP BEGIN at /home/jab29/scratch-midway2/project-2/src/nbody_cpu_serial.c(21,3)
remark #15542: loop was not vectorized: inner loop was already vectorized
```

```
LOOP BEGIN at /home/jab29/scratch-midway2/project-2/src/nbody_cpu_serial.c(24,5)
remark #15300: LOOP WAS VECTORIZED
LOOP END
```

```
LOOP BEGIN at /home/jab29/scratch-midway2/project-2/src/nbody_cpu_serial.c(24,5)
<Remainder loop for vectorization>
LOOP END
LOOP END
```

Report from: Code generation optimizations [cg]

/home/jab29/scratch-midway2/project-2/src/nbody\_cpu\_serial.c(25,18):remark #34032: adjacent

## 10

### L1/L2/L3 Cache Hit Rate CPU

~For running with n=50000 (need to finish the assignment)~ Note: For running with n=10000 (witnessed weird behavior w/50000)

We have recorded cache hit rates in the following table...

Cache Hit Rate									
CPU code	L1 Misses	Total Loads and Stores*	L1 Hit Rate	L2 Misses	L2 Accesses	L2 Hit Rate	L3 Misses	L3 Accesses	L3 Hit Rate
Serial	749487965	12003083184	0.937558721079342	247001975	747804274	0.669697026898779	22	394470262	0.999999944229002
Multicore	903058336	16739712256	0.946052935546947	508274944	899674560	0.435045774774381	11996032	347276768	0.96545685428632

A couple notes:

- L1 cache rate stays high for both implementations. I imagine this comes from data being contiguous and when L1 lines come in we simply automatically benefit by each thread operating on contiguous elements.
- L2 and L3 rates suffer for the multithreaded implementation, this is within expectations especially considering that there is quite the high amount of false-sharing going on when each thread has to visit the entire universe of bodies.

## 11

### Branch misprediction rate

For running with n=50000 (need to finish the assignment)

PAPI\_BR\_MSP = 1006077

PAPI\_BR\_INS = 100002253817

Mispredict rate = 0.001006054 %

The misprediction rate is small for this code since the majority of branching lives in looping code. These branches can easily be predicted with the simplest of prediction methods, hence the small mispredict rate.

## 12

### TLB Hit Ratio

For running with n=50000 (need to finish the assignment)

PAPI\_TLB\_DM = 20619

Note: I can't find an appropriate counter to use as a baseline for the TLB hit rate. I only have total misses. :(

Now in general, I don't expect many TLB misses in our code. Most of our code operates on the same contiguous elements of memory, which should fit nicely with regards to address mappings. Though our code is  $O(n^2)$  our memory is still  $O(n)$ .



In general, TLB performance shouldn't have much of an impact on code performance unless you have many nbodies, where you start jumping past the sizes of the addresses managed in your TLB.

## 13

### Hardware Thread Gains

Our lscpu indicates that there is 1 thread per CPU. So hyperthreading isn't enabled yet. I'm not sure how to "enable" this at a system level so all I can do is speculate.

I imagine enabling hardware threading (hyperthreading) would add to some performance improvements just given the compute nature of this code.

## 14

### Thread count memory saturation

Memory Util					
Num Threads	L3 Cache Misses	Across All Threads	Exec Time	Throughput	*actual
8	2953369	23626952	138.759662	170272.481638071	170272.481638071
16	2144277	34308432	70.649088	485617.478883804	485617.478883804
32	1499496	47983872	40.649232	1180437.35734048	1180437.35734048
64	1064815	68148160	40.634175	1677114.39939411	838557.199697053
128	954084	122122752	40.628149	3005865.51457218	751466.378643044
256	993872	254431232	40.657920	6257851.65596273	782231.456995341
512	1039198	532069376	40.624539	13097240.9557681	818577.559735509

Figure 3: Memory Util

The memory never gets fully saturated. For this measurement, I couldn't find an appropriate memory activity measure, so I used L3 misses (which should yield a memory access) as an appropriate proxy.

This makes sense though, the memory bandwidth limit for this system is quite high. Additionally, the false-sharing of bodies in the inner loop shouldn't outpace this bound.

Note: The "actual" column refers to the throughput assuming 32 active threads at the same time (since anything higher is bounded by the CPU)

## 15

### Max parallel efficiency

(Measurements in section below).

We hit 32 threads for maximum parallel efficiency. This is illustrated quite so in our strong-scaling, which hits an exact wall-time of 40s for any number of threads past 32.

This is explained since there are 28 cores that are online on our system. We haven't had any indication that we reached a memory bound before this limit.

## 16

### Strong-scaling/Weak-scaling

#### Strong Scaling

#### Weak Scaling

Note: Our weak-scaling assumes a linear input increase when in reality this refers to a quadratic input increase per thread (due to  $O(n^2)$  complexity of the algorithm. So the reported as-is weak scaling might be frightening without this context. We've only reported the parallel efficiency (scaled to threads) because of the complications.

We can see that we saturate performance at 32 threads. This comes from the processor having 28 on-line CPUs (running for 28 threads yields the same performance). Weak scaling also shows us that further scaling past 32 leads to unsustainable runtimes.

## 17

### L1 miss rate serial v. multicore

Based upon our results in section 10, L1 cache hit rates are roughly the same for both implementations. I think this mostly has to do with the iteration in general being on contiguous sections of memory, and shouldn't be affected too much in a multithreaded environment. A given thread will operate on all other bodies (and thus will suffer from false sharing), but when those cache lines are loaded the multicore implementation should benefit from the next set of accesses to contiguous memory.

The above is different from L2/L3 cache where most of the time a thread will attempt to look in those caches (usually non contiguously), there's a relatively high probability that its already been modified.




Figure 4: StrongScaling




Figure 5: WeakScaling

## 18

### GPU Optimization Explanation

#### gpu1

This implementation tries to parallelize computation with both CPU threads as well as GPU threads. CPU for the post-calc position update and GPU for the inter-force calculation.

The CPU threading is straightforward, parallelize the simple (honestly SIMD) update.

The GPU side is broken up across 1 axis of bodies, with each thread responsible for at most 1 body.

#### gpu2

Similar to GPU1 but without openmp for the position update. Probably better considering the relative cost of position update.

#### gpu3

For the GPU block, write the starting position state for all bodies first as to avoid possible contention between threads.

My initial guess is that this would help from false-sharing perspective (not to mention, correctness?)

The syncthreads is there to converge on a state for each “subblock” of our routine.

EDIT: After refreshing my mind a bit, shared memory should be faster to access than the default global memory, hence the reason for this optimization.

#### gpu4

This adds a loop unroll for our loop per gpu thread. A probable way to improve occupancy by shedding away temporary state variables.

## 19

### Fraction of host/device copy

Host/Device Copy					
Version	N	mem_ops_ms	body_calc_ms	Percentage	
GPU1	1000	205.59	6.638	96.8722317507586	
GPU1	10000	281.182	84.025	76.9925001437541	
GPU1	100000	6912.30	6893.86	50.0667817843629	
GPU1	500000	164682	164583	50.0150334836682	
GPU2	1000	191.49	4.0751	97.9162437469671	
GPU2	10000	244.86	53.300	82.1236919774618	
GPU2	100000	4033.03	4010.60	50.1394270994563	
GPU2	500000	95646.0	95520.1	50.0329294786053	
GPU3	1000	191.32	2.2146	98.8557084883013	
GPU3	10000	227.976	39.024	85.3842696629213	
GPU3	100000	3221.67	3199.28	50.1743511474159	
GPU3	500000	76294.6	76166.4	50.0420435390034	
GPU4	1000	189.1071	2.5783	98.6549314658289	
GPU4	10000	243.263	40.084	85.8533882483316	
GPU4	100000	3227.57	3205.46	50.1718474808916	
GPU4	500000	76637.0	76526.1	50.0362032369415	

It makes sense that the cuda memcpy/malloc operations would be amortized for larger inputs. What's odd to me is this 50% barrier. It got these values from nvprof but I am somewhat confused on why this is the case. Not sure if I mistook the api calls section to not actually include the gpu computation as well. 50% asymptote seems odd to say the least.

## 20

### GPU Occupancy

I had trouble trying to get this to work on nvprof. All I can do here is speculate (and assume that a quicker runtime is due to higher occupancy, which shouldn't be a stretch in GPU-land)

We arrive at higher occupancy in each of our optimizations as follows.

- From 2 to 3, since we take advantage of shared memory, each thread's access to the universe of bodies will lead to less stalling. This will lead to warps executing more deterministically which should lead to higher occupancy.
- From 3 to 4, since we perform loop unrolling, we can remove register use for temporary variables. This should lead to less contention among warps and lead to higher occupancy.

## Addendum

### Resource Stats

#### CPU

jab29@midway2-0121

Architecture: x86\_64  
CPU op-mode(s): 32-bit, 64-bit  
Byte Order: Little Endian  
CPU(s): 56  
On-line CPU(s) list: 0-27  
Off-line CPU(s) list: 28-55  
Thread(s) per core: 1  
Core(s) per socket: 14  
Socket(s): 2  
NUMA node(s): 2  
Vendor ID: GenuineIntel  
CPU family: 6  
Model: 79  
Model name: Intel(R) Xeon(R) CPU E5-2680 v4 @ 2.40GHz  
Stepping: 1  
CPU MHz: 1200.000  
CPU max MHz: 2401.0000  
CPU min MHz: 1200.0000  
BogoMIPS: 4799.81  
Virtualization: VT-x  
L1d cache: 32K  
L1i cache: 32K  
L2 cache: 256K  
L3 cache: 35840K  
NUMA node0 CPU(s): 0-13  
NUMA node1 CPU(s): 14-27  
Flags: fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse

#### Mem/Cache

### Memory Cache and TLB Hierarchy Information.

TLB Information.

There may be multiple descriptors for each level of TLB if multiple page sizes are supported.

L1 Data TLB:  
Page Size: 4 KB  
Number of Entries: 64  
Associativity: 4

Cache Information.

L1 Data Cache:

Total size:	32 KB
Line size:	64 B
Number of Lines:	512
Associativity:	8

L1 Instruction Cache:

Total size:	32 KB
Line size:	64 B
Number of Lines:	512
Associativity:	8

L2 Unified Cache:

Total size:	256 KB
Line size:	64 B
Number of Lines:	4096
Associativity:	8

L3 Unified Cache:

Total size:	35840 KB
Line size:	64 B
Number of Lines:	573440
Associativity:	20

mem\_info.c

PASSED